



MICROSERVICES

SERVICE DISCOVERY & API GATEWAYS





CONTENT

- Protocols
- Service Discovery
- API Gateways



CONTENT - PROTOCOLS

- Gossip
- Raft



CONTENT - SERVICE DISCOVERY

- Client-side service discovery
- Server-side service discovery
- DNS
- Key-value stores



CONTENT - API GATEWAYS

- Basics
- Load balancing strategies
- Failover techniques



PROTOCOLS



PROTOCOLS - GOSSIP

- Protocol to:
 - keep a cluster state in sync
 - manage the clusters health by constantly checking which nodes are available
- Used by Consul (Serf) based on **Scalable Weakly-consistent Infection-style Process Group Membership Protocol (SWIM)**


GOSSIP - PRINCIPLES - PART 1

- Periodic, pairwise, inter-process (or network) interactions
- The information exchanged during these interactions is of bounded size
- Agents are synchronizing their state when they interact with each other
- Reliable communication is not assumed


GOSSIP - PRINCIPLES - PART 2


- The frequency of the interactions is low compared to typical message latencies so that the protocol costs are negligible.
- There is some form of randomness in the peer selection. Peers might be selected from the full set of nodes or from a smaller set of neighbors.
- Due to the replication there is an implicit redundancy of the delivered information.


Source




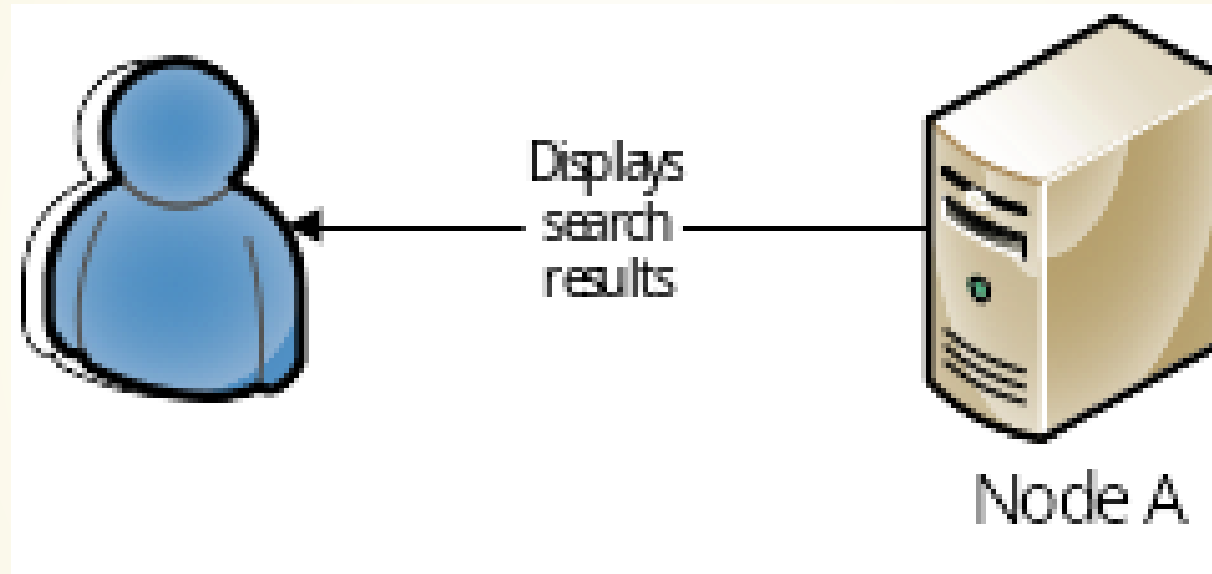
(c) Gossip-based approach,
where peers operate in parallel, and
each peer communicates with one or
more randomly selected partner











GOSSIP SEARCH SAMPLE - REMARKS

- A search query should "age out" after a given time to reduce traffic
- If there are many search queries a maximum of data that may be exchanged during one "gossip" has to be defined
- Given a frequency of 10 gossips per second, a maximum of 30 rounds of gossip per search query and a network of 25.000 machines a query would take just about 3 seconds!



PROTOCOLS - RAFT CONSENSUS

RAFT CONSENSUS BASICS

- Designed as an easier alternative to Paxos
- Uses leader election to achieve consensus
- Models a distributed state machine
 - Every node is a state machine
 - All nodes have to apply the same commands in the same order to stay in sync (same resulting state/transition)

RAFT CONSENSUS BASICS

- Just one leader in a Raft cluster, all other nodes are followers
- Leader is responsible for the log replication to all followers
- Followers are expecting a heartbeat within a given timeout otherwise they suspect the leader failing
- If a leader fails a new leader is elected

Visualization <https://raft.github.io/>



RAFT - LEADER ELECTION

- Leader election is started by a candidate server (a server that wasn't contacted by the leader within the timeout period)
- Candidate increments the term number (serial for periods where a leader was elected) and proposes itself as the new leader and sends a message to all other servers requesting their vote



RAFT - LEADER ELECTION

- If candidate gets a response with a term number at least as large as his current term number the election is defeated and the candidate is switching in follower mode
- If the candidate server gets a majority of votes he's getting the new leader
- If neither happens (split vote) a new term is getting started (resulting in a new election)

RAFT - LOG REPLICATION

- Leader replicates received requests (commands for the state machine) to all followers
- Leader appends the command to his log as a new entry and sends a `AppendEntry` to the followers
- When the leader receives confirmation of a majority of his followers he applies the entry to his state machine (request is considered committed)

RAFT - LOG REPLICATION

- When a follower learns that an entry was applied by the leader he applies the entry to his local state machine
- In case of a leader crash the new leader enforces a replication of his log to all followers. To get a consistent state the leader compares his log with every log of the followers, takes the latest where they agree and replaces all following entries with his own

RAFT - SAFETY RULES

- Election safety (at most one leader can be elected in a given term)
- Leader Append-Only (a leader can only append new entries to its logs - it can neither overwrite nor delete entries)
- Log Matching (if two logs contain an entry with the same index and term, then the logs are identical in all entries up through the given index)

RAFT - SAFETY RULES

- Leader Completeness (if a log entry is committed in a given term then it will be present in the logs of the leaders since this term)
- State Machine Safety (if a server has applied a particular log entry to its state machine, then no other server may apply a different command for the same log)

SERVICE DISCOVERY BASICS

Service discovery mechanisms are required for multiple tasks in a microservice environment:

- server resolution for cross service communication
- dynamic load balancer configuration
- dynamic monitoring configuration
- ...



APPROACHES

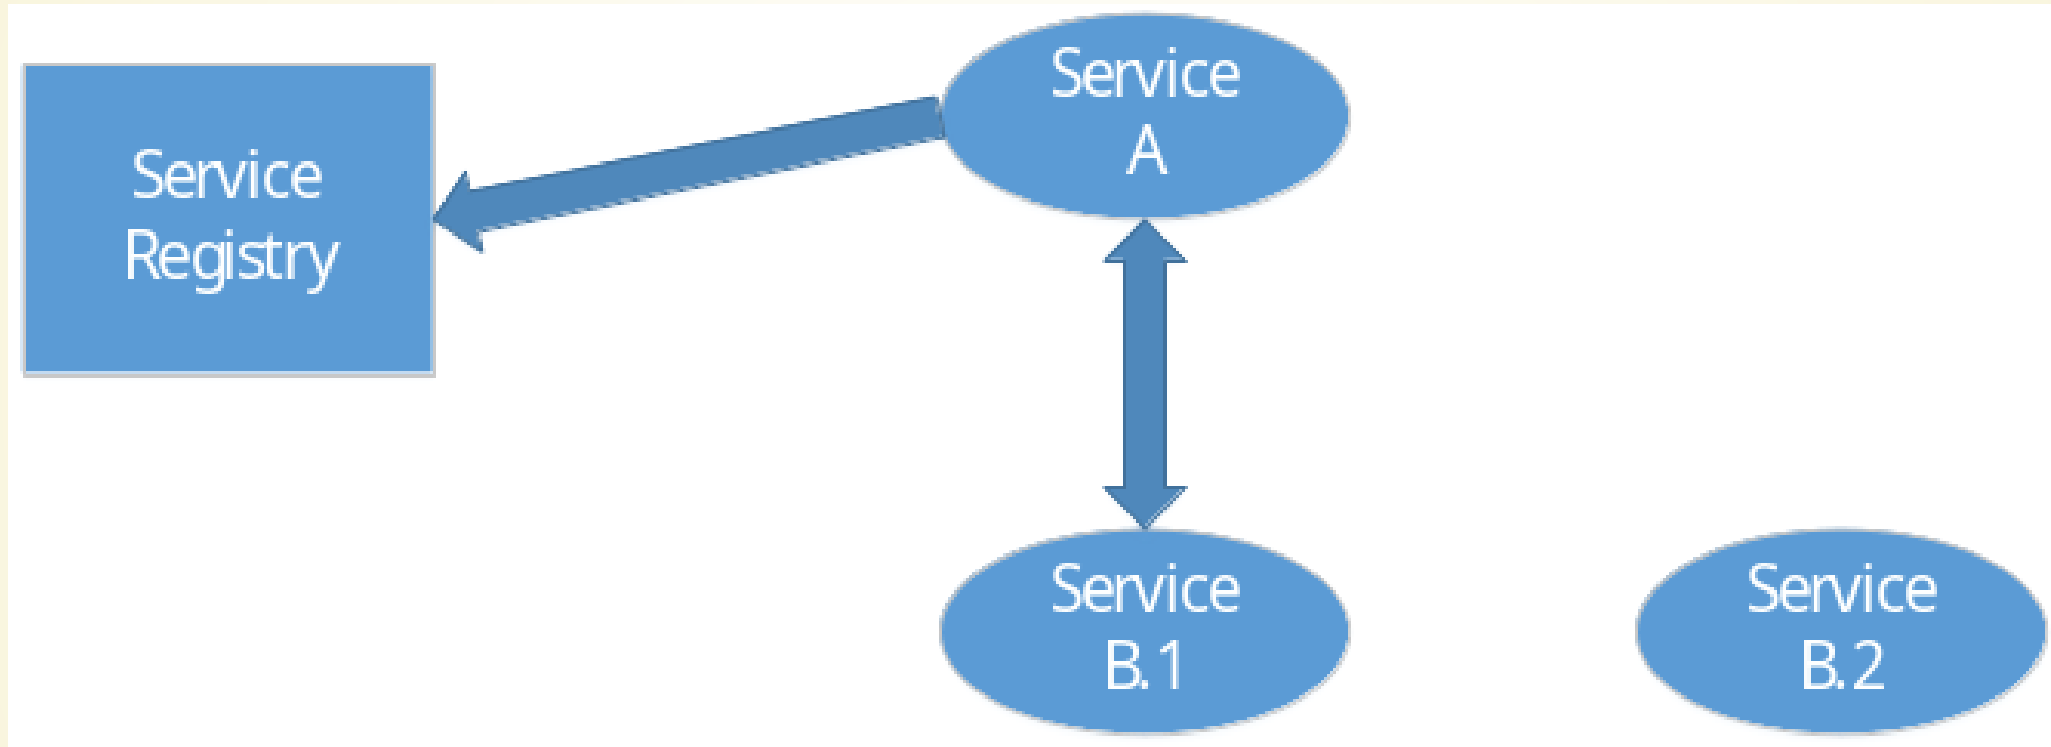
- Static configuration files (for the sake of completeness...)
- DNS based solutions (e.g. in Kubernetes with CoreDNS)
- Specialized products e.g. Eureka & Consul



CLIENT-SIDE VS. SERVER-SIDE SERVICE DISCOVERY



CLIENT-SIDE SERVICE DISCOVERY






WORKFLOW

1. Service A/B.1/B.2 starts and registers itself at central service registry
2. Service A asks for one endpoint/all endpoints of Service B
3. Service Registry returns one/all endpoints to Service A
4. Depending on the implementation Service A decides which instance to call or calls the only it is aware of



SERVER-SIDE SERVICE DISCOVERY



WORKFLOW

1. Service A/B starts and registers itself at central service registry
2. Service A sends request to central router (e.g. an API Gateway)
3. Gateway is redirecting call to concrete endpoint like a proxy
4. Gateway returns the response it got from the concrete endpoint to the original caller



SERVICE REGISTRATION





SERVICE REGISTRATION - SELF REGISTRATION

- Every client/service registers himself at the service registry
- Every client has to deregister himself on failures or when quitting
- Every client has to deal with the API of the service registry himself
- E.g. Netflix Eureka

SERVICE REGISTRATION - 3RD PARTY REGISTRATION

- Clients/services are registered by a external instance
- Whenever a client exits the external component deregisters him\
- The external component has to monitor every known service to ensure that it's still available
- E.g. registrator, Nomad



DNS BASED SERVICE DISCOVERY

DNS - VIABLE RECORD TYPES - PART 1

Record name	Explanation
-------------	-------------

A or AAAA	Host entries (e.g. www.google.de – IPv4: 172.217.21.35 and IPv6:2a00:1450:4016:80d::2003)
-----------	---

CNAME	Alias of a host entry (e.g. www.fh-rosenheim.de and fh-rosenheim.de)
-------	--

SRV	Service location record (includes port of the service)
-----	--

DNS - VIABLE RECORD TYPES - PART 2S

Record name	Explanation
-------------	-------------

TXT	Often carries machine-readable data (often used e.g. for domain validation in Azure, C&C servers,...)
------------	---

NAPTR	Name Authority Pointer – allows regular-expression-based rewriting of domain names (e.g. to form URIs)
--------------	--

Source

DNS AS SERVICE REGISTRY

- A (or AAAA) can be used to locate services (a single A record may contain multiple IP addresses e.g. amazon.com)
- SRV records are even better because it's also possible to store the port of the service in a SRV record
- Every instance has to register itself at a DNS server or a 3 rd party service has to look for new instances and register them within a DNS server
- Developers and administrators are required to create a common schema for service naming

DNS - NAMING SCHEMAS

Schema sample

Use case

`<servicename>-
<env>.domain.tld`

All environments share
the same domain/DNS
server

`<servicename>.
<env>.domain.tld`

Subdomain per
environment (e.g.
test.domain.tld and
staging.domain.tld,
keep prod on
domain.tld)

DNS AS SERVICE REGISTRY - CONSIDERATIONS

- Relatively easy to implement
- No special software/libraries required
- TTL of entries might lead to stale entries
- DNS caching
- Requires special DNS server implementation to support dynamic registration

KEY-VALUE STORES FOR SERVICE DISCOVERY

Classical ones:

- ZooKeeper
- etcd

Specific for Microservices:

- Consul
- Eureka

CLASSICAL KEY-VALUE STORES

- Developed as distributed configuration stores
- Hierarchical structured
- Normally offer some kind of "watches" or "subscriptions"

MICROSERVICE SPECIFIC ONES

- Implement specific domain knowledge (special entities for services and endpoints)
- Offer possibilities to register e.g. health checks to ensure that registered services are available
- Some are also offering configuration stores (e.g. Consul)
- Various APIs (e.g. HTTP or DNS)



LOAD BALANCING BASICS


- Load Balancers are needed to avoid single point of failures
- Load Balancers distribute calls sent to them to one or more instances
- Load Balancers keep track of their known backends to avoid errors when a services is no longer healthy

LOAD BALANCING BASICS

- Optionally they have additional features like SSL termination
 - Only one or a few servers where certificates have to be exchanged when a new certificate is created
 - No special handling required in the services behind the load balancer
 - Admins have to take care that communication between load balancer(s) and nodes is safe (e.g. VLANs)

LOAD BALANCING BASICS

- In Microservice environments it's essential that the load balancer(s) can be reconfigured dynamically (e.g. in combination with etcd or Consul)
- When you're able to scale your microservice instances but not the persistence layer of them you're just moving the single point of failure one layer backwards!





LOAD BALANCING STRATEGIES

- Round-robin
- Weighted round-robin
- Least connection
- Weighted least connection
- Agent Based Adaptive Load Balancing
- Chained Failover (Fixed Weighted)
- Weighted Response Time
- Source IP Hash
- ...



(WEIGHTED) ROUND-ROBIN

- There are many Round-robin algorithms
- The simplest one is to use a FIFO queue to keep track of all available backends
 1. Dequeue
 2. Relay request
 3. Enqueue

(WEIGHTED) ROUND-ROBIN

- Results in max-min-fairness (the longest waiting requests gets the highest priority)
- The weighted round robin algorithm gives every backend a weight and the scheduler takes the weights into account to prefer servers with a higher weight before servers with a lower weight (e.g. used for quality of service (QoS))



(WEIGHTED) LEAST CONNECTION

- In contrast to round-robin the least connection algorithm takes the load of every node in account.
- The least connection algorithm relays an incoming request always to the node which has the lowest count of active connections.



(WEIGHTED) LEAST CONNECTION

- This way nodes with a higher performance handle more requests than nodes with lower performance without the need to configure weights.
- The weighted least connection variant enables the administrator to give nodes a weight. These weights are considered when two nodes serve the same count of active connections and the node with the higher weight is considered first.



AGENT BASED ADAPTIVE LOAD BALANCING

- Every node has an local agent installed which reports real time data to the load balancer (e.g. CPU usage, memory allocation,...)
- Load balancers takes load of every node into to account when a new request has to be relayed
- Can be combined with weighted round-robin or weighted least connection algorithms
- Used e.g. in Windows Terminal Server (RDS role after Server 2003)



CHAINED FAILOVER

- All backend nodes are in a predefined chain
- Whenever the first node can't handle/accept another request the next node in the chain is taken into account and so on
- Not a real load balancing protocol!

WEIGHTED RESPONSE TIME

- Kind of health check done by the load balancer(s)
- Uses the response time of the health check to determine the fastest server currently available
- Whenever a node is under heavy load the response times will be longer than the response times of a node with least load.
- Avoid overload of nodes.



SOURCE IP HASH

- Algorithm creates a hash of source and destination IP (unique hash key)
- Hash key is used to determine to which node the request should be forwarded

SOURCE IP HASH

- When the same client sends another request the hash key can be regenerated and the client gets forwarded to the same node
- Useful for stateful services (don't do that in microservices!) when nodes aren't able to sync session information because a client always gets relayed to the same node (as long as its source IP does not change)



NETWORK FAILOVER

- But wait! If I have 1 load balancer what happens if this load balancer fails?
- Possible solution: multiple load balancers with multiple DNS A/AAAA records to balance the load of the load balancers -> but DNS does not check for availability and there are the caches...

NETWORK FAILOVER

- Better solution: configure network based failover:
 - Common address redundancy protocol (CARP)
 - Gateway Load Balancing Protocol (GLBP) (just for routers)

COMMON ADDRESS REDUNDANCY PROTOCOL (CARP)

- Enables multiple hosts in the same LAN to share a set of IP addresses
- Available on BSD and Linux based hosts
- Master-slave (or more polite active-passive) based
- One master per group of redundancy
- Each group of redundancy shares one virtual IP
- A server maybe member of multiple groups of redundancy



COMMON ADDRESS REDUNDANCY PROTOCOL (CARP)

GATEWAY LOAD BALANCING PROTOCOL (GLBP)

- Proprietary protocol created by Cisco for redundant routers
- Allows weighting parameter to be set
- Based on the weights (in a virtual router group) ARP requests will be answered with MAC addresses pointing to different routes
- Balances in round-robin fashion by default



GATEWAY LOAD BALANCING PROTOCOL (GLBP)

- Elects one Active Virtual Gateway (AVG) for each group
- AVG assigns every listener (and itself) virtual MAC addresses which enables Active Virtual Forwarders (AVF)
- Each AVF is responsible to forward packages sent to its virtual MAC address

API GATEWAYS


- API Gateways are a crucial part of every microservices environment
- API Gateways enable a microservices environment to scale by implement load balancing (e.g. round-robin based)
- Access to specific services is managed by:
 - DNS-Host per Service (A/AAAA/CName e.g. ServiceA.my-domain.com)
 - Virtual Routes (e.g. gateway.my-domain.com/ServiceA)

API GATEWAYS

- API Gateways are the single entry point for your whole application
- API Gateways are a kind of server-side service discovery (usually just for client apps but it's possible to use it also for cross service calls)

API GATEWAYS

- They also hide implementation details by optionally aggregating all internal APIs to one (or more in case of Backends for frontends) in the point of view of the client app(s)
- In the case of custom API Gateways the gateway may also execute calls to multiple services and aggregate the responses answering to the client request



BACKEND FOR FRONTENDS

- Configure a API Gateway per kind of frontend e.g.
 - One for your web app
 - One for your mobile app
 - One for all 3 rd party applications (public API)
- The backends for frontends-pattern ensures the optimal API for every kind of application (e.g. gRPC gateway for desktop apps but RESTful API for web apps)

